

# Lecture 1: Greedy Algorithms

Yury Makarychev

# Administrative

- Lectures: on Tuesdays and Thursdays, 9:30-10:50am
- Tutorials: on Wednesdays, 4:30-5:20pm
- TA: Theodoros Papamakarios
- Grader: Francisco Mendes
- Textbook: *Algorithm Design* by Kleinberg and Tardos
- Please make sure that you have access to the [Canvas](#) webpage for the course!

# Assignments

4 homework assignments (60%) + 1 final exam (40%)

Each HW will have 3 problems + 1 programming assignment

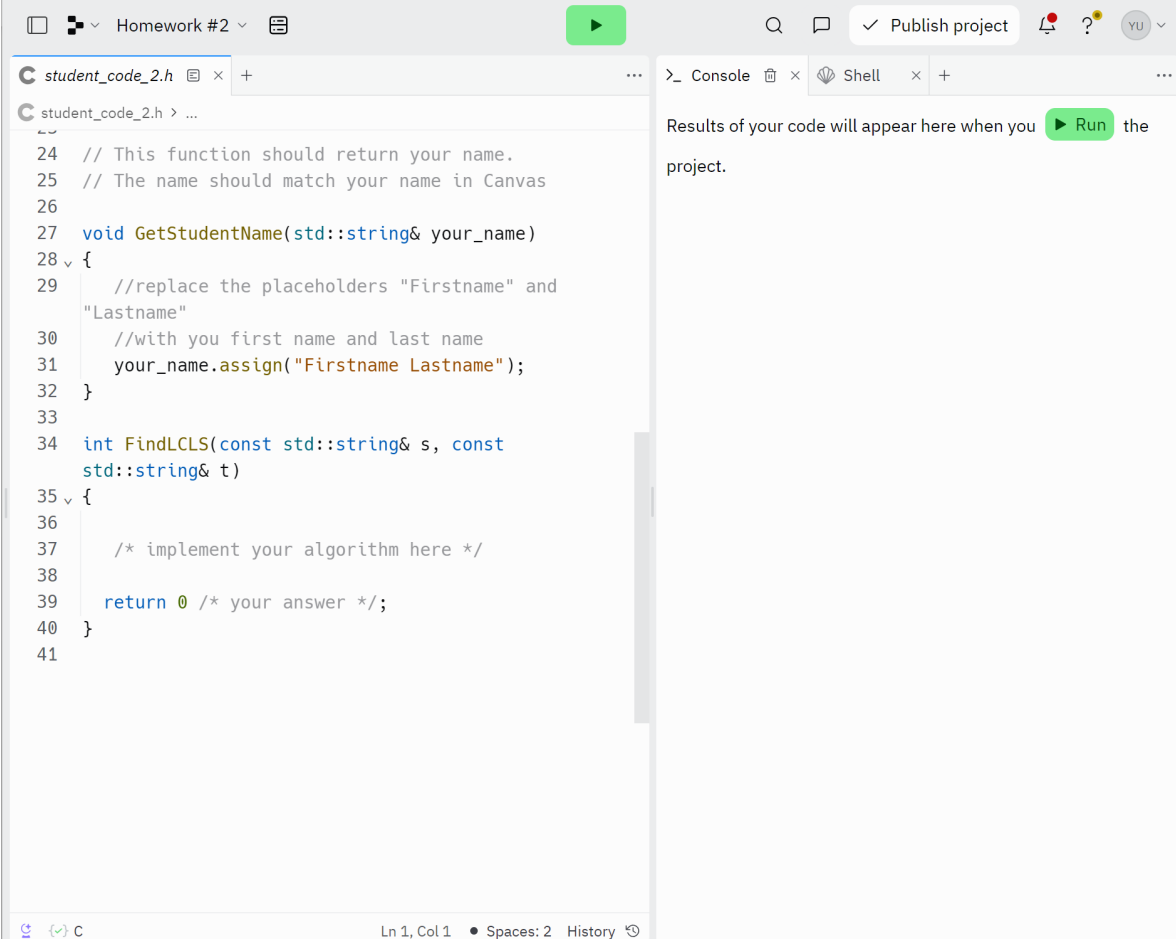
Please submit your text solutions via [Gradescope](#) and programming assignment solutions on [Canvas](#).

- First HW will be posted next Thursday.

# Assignments

## Programming assignment:

- Need to design and implement an algorithm in C++.
- Use [replit.com](https://replit.com). You will need to implement only one function.



The screenshot shows a Replit IDE interface for a C++ project titled "Homework #2". The main editor window displays the file `student_code_2.h` with the following code:

```
24 // This function should return your name.  
25 // The name should match your name in Canvas  
26  
27 void GetStudentName(std::string& your_name)  
28 {  
29     //replace the placeholders "Firstname" and  
30     //with you first name and last name  
31     your_name.assign("Firstname Lastname");  
32 }  
33  
34 int FindLCLS(const std::string& s, const  
35 std::string& t)  
36 {  
37     /* implement your algorithm here */  
38  
39     return 0 /* your answer */;  
40 }  
41
```

On the right side, there is a "Console" and "Shell" tab. The console area contains the text: "Results of your code will appear here when you [Run](#) the project." with a green "Run" button next to it. The top bar includes a "Publish project" button and a user profile icon labeled "YU". The bottom status bar shows "Ln 1, Col 1", "Spaces: 2", and a "History" link.

# Questions?


# Syllabus

- Greedy Algorithms
- Dynamic Programming
- Max Flow, Min Cut, and Their Applications
- Linear Programming
- Classes P and NP. NP-hardness.
- Approximation Algorithms
- Multiplicative Weight Updates
- (if we have any time left) more advanced topics

# Greedy Algorithms

# Job (Interval) Scheduling Problem

- Given a set of intervals  $I_1 = (s_1, t_1), \dots, I_n = (s_n, t_n)$  on the real line.




- Interval  $I_i$  represents a job that starts at  $s_i$  and finishes at  $t_i$
- Need to schedule a subset of jobs on a machine/computer that can execute only one job at a time
- Two jobs  $I_i$  and  $I_j$  are *compatible* if they don't overlap:  
either  $t_i \leq s_j$  or  $t_j \leq s_i$




# Job Scheduling Problem: Examples

A *feasible* solution/schedule is a subset of jobs  $S \subseteq \{I_1, \dots, I_n\}$  such that every two jobs in  $S$  are compatible.

$\{I_1, I_2, I_3\}$  is a feasible solution:




$\{I_1, I_3, I_4\}$  is **not** a feasible solution, because  $I_1$  and  $I_4$  overlap:



# Job Scheduling Problem

- Given a set of jobs  $I_1 = (s_1, t_1), \dots, I_n = (s_n, t_n)$ .
- Find a feasible solution  $S$  of maximal size.


I.e., find the largest number of intervals no two of which overlap.



# Greedy Algorithm

We will design a *greedy algorithm*.

A greedy algorithm constructs an optimal solution step-by-step. At each step, it makes a choice that is “locally optimal”.




# Greedy Algorithm

We will design a *greedy algorithm*.

A greedy algorithm constructs an optimal solution step-by-step. At each step, it makes a choice that is “locally optimal”.

- $S = \emptyset$
- while there is a job that starts after all jobs in  $S$  finish
  - among all such jobs, find a job  $I_j$  with the least value of  $t_j$  (that finishes first)
  - add  $I_j$  to  $S$




# TO DO items

- Prove that the algorithm finds a **feasible** solution.
- Prove that the algorithm finds an **optimal** solution:

$$|S| \geq |S'| \text{ for every feasible solution } S'$$

- Discuss how to implement the algorithm efficiently and find its running time.

# Feasibility



- $S = \emptyset$
- while there is a job that starts after all jobs in  $S$  finish
  - among all such jobs, find a job  $I_j$  with the least value of  $t_j$  (that finishes first)
  - add  $I_j$  to  $S$

Why is the set of jobs  $S$  returned by the algorithm a feasible solution?

**Proof:** Consider two jobs  $I_i$  and  $I_j$  in  $S$ .

Assume that  $I_i$  was added to  $S$  before  $I_j$ .

When we added  $I_j$ , it was compatible with all jobs in  $S$ , including  $I_i$ . ■


# Optimality

Consider an optimal solution  $S^*$ . Prove that  $|S| = |S^*|$ .

Proof:

Sort jobs in  $S$  and  $S^*$  from left to right. Let


$$S = \{I_{i_1}, \dots, I_{i_k}\} \text{ and } S^* = \{I_{j_1}, \dots, I_{j_{k^*}}\}$$



Note  $k \leq k^*$ , since  $S^*$  is an optimal solution. Want:  $k \geq k^*$ .

# Optimality

Prove by induction on  $\ell$  that  $t_{i_\ell} \leq t_{j_\ell}$  for  $\ell \in \{1, \dots, k\}$ .




Base case:  $t_{i_1} \leq t_{j_1}$ . Why?



# Optimality

Prove by induction on  $\ell$  that  $t_{i_\ell} \leq t_{j_\ell}$  for  $\ell \in \{1, \dots, k\}$ .



Induction step: Assume that  $t_{i_{\ell-1}} \leq t_{j_{\ell-1}}$ .

- $I_{j_\ell}$  lies to the right of  $I_{i_{\ell-1}}$ :  $s_{j_\ell} \geq t_{j_{\ell-1}} \geq t_{i_{\ell-1}}$ .
- When we added  $I_{i_\ell}$  to  $S$ ,  $I_{j_\ell}$  was to the right of  $I_{i_{\ell-1}}$ . We could have added  $I_{j_\ell}$  but instead added  $I_{i_\ell}$ . Why?

# Optimality

Prove by induction on  $\ell$  that  $t_{i_\ell} \leq t_{j_\ell}$  for  $\ell \in \{1, \dots, k\}$ .

Induction step: Assume that  $t_{i_{\ell-1}} \leq t_{j_{\ell-1}}$ .

- $I_{j_\ell}$  lies to the right of  $I_{i_{\ell-1}}$ :  $s_{j_\ell} \geq t_{j_{\ell-1}} \geq t_{i_{\ell-1}}$ .
- When we added  $I_{i_\ell}$  to  $S$ ,  $I_{j_\ell}$  **was to the right** of  $I_{i_{\ell-1}}$ . We could have added  $I_{j_\ell}$  but instead added  $I_{i_\ell}$ . **Why?**
- It had to be the case that:  $t_{i_\ell} \leq t_{j_\ell}$ .
- We proved  $t_{i_\ell} \leq t_{j_\ell}$  for  $\ell \in \{1, \dots, k\}$ . ■

# Optimality

Proved by induction that  $t_{i_\ell} \leq t_{j_\ell}$  for  $\ell \in \{1, \dots, k\}$ .

Remains to show that  $k \geq k^*$ . Assume to the contrary that  $k < k^*$ .

- $S = \emptyset$
- while there is a job that starts after all jobs in  $S$  finish
  - ...

The algorithm terminates after adding  $I_{i_k}$  to  $S$ .

Thus, there is no job to the right of  $I_{i_k}$  in  $S$ . However,  $s_{j_{k+1}} \geq t_{j_k} \geq t_{i_k}$ .

Contradiction. We proved that  $S$  is an optimal solution.

# Implementation

- Sort all jobs by their stop time  $t_i$ :  
$$t_1 \leq t_2 \leq \dots \leq t_n$$
- Add  $I_1$  to  $S$
- $T = t_1$      (*the termination time of the last currently scheduled job*)
- $i = 2$      (*the first unprocessed job*)
- while  $i \leq n$  do
  - if  $s_i \geq T$ , then
    - add  $I_i$  to our schedule  $S$
    - $T = t_i$
  - $i = i + 1$

# Implementation

Running time

- Sort all jobs by their stop time  $t_i$ :

$$t_1 \leq t_2 \leq \dots \leq t_n$$

$O(n \log n)$

- Add  $I_1$  to  $S$
- $T = t_1$  (the termination time of the last currently scheduled job)
- $i = 2$  (the first unprocessed job)
- while  $i \leq n$  do
  - if  $s_i \geq T$ , then
    - add  $I_i$  to our schedule  $S$
    - $T = t_i$
  - $i = i + 1$

$O(n)$

# Combinatorial Optimization

Job Scheduling is a combinatorial optimization problem.

In a combinatorial optimization problem:

- We are given an instance of the problem, which defines the set of feasible solutions and an objective function.
- The goal is to find an optimal solution, a feasible solution whose objective is
  - greater than or equal to (for a maximization problem), or
  - smaller than or equal to (for a minimization problem)

↑ value

↓ cost

than the value of any other feasible solution.

# Combinatorial Optimization

Job Scheduling is an example of a large class of scheduling problem. It's also known as [Maximum Independent Set in Interval Graphs](#).

In this class, we will mostly study various approaches – such as Greedy Algorithms, Dynamic Programming, and Linear Programming – for solving combinatorial optimization problems.

Today, we designed a greedy algorithm for Job Scheduling. We used a “stay ahead” type argument: our solution always “stays ahead” of an optimal solution ( $t_{i_\ell} \leq t_{j_\ell}$ ).

# Greedy Algorithms: Pros and Cons

- + Very efficient
- + Easy to implement and analyze
- + If there is a greedy algorithm for a problem, use it!
- There are no greedy algorithms for many problems
- Even if there is a greedy algorithm for a problem, there may be no algorithm for its variant (e.g. for Weighted Job Scheduling).




# Minimum Weighted Completion Time

➤ Given  $n$  jobs.

Each job  $i$  has duration  $t_i$  and weight  $w_i$ .

Jobs don't have fixed start and stop times.

➤ We want to run all  $n$  jobs one after another:




It's up to us to decide the order in which we run the jobs.

# Minimum Weighted Completion Time

Given an order  $i_1, \dots, i_n$ , we schedule the jobs as follows:

- job  $i_1$  starts at time 0 and completes at  $t_{i_1}$
- job  $i_2$  starts at time  $t_{i_1}$  and completes at  $t_{i_1} + t_{i_2}$
- ...



Completion time of job  $i_\ell$  is

$$c(i_\ell) = \sum_{a=1}^{\ell} t_{i_a}$$

# Objective

$$c(i_\ell) = \sum_{a=1}^{\ell} t_{i_a}$$

Note that  $c(i)$  depends on the ordering  $i_1, \dots, i_\ell$ .

- Find an ordering  $i_1, \dots, i_n$  that minimizes the weighted sum of completion times:



$$\sum_{i=1}^n w_i c(i)$$

# Warm Up

$$c(i_\ell) = \sum_{a=1}^{\ell} t_{i_a}$$

Consider a couple of examples. Let  $n = 2$ .

1. Assume that  $t_1 = t_2 = 1$ . There are two solutions:



Their costs are

$$w_1 + 2w_2$$

and

$$w_2 + 2w_1$$



It's better to put **heavier** jobs first.

# Warm Up

$$c(i_\ell) = \sum_{a=1}^{\ell} t_{i_a}$$

Consider a couple of examples. Let  $n = 2$ .

2. Assume that  $w_1 = w_2 = 1$ . There are two solutions:



Their costs are

$$2t_1 + t_2$$

and

$$t_1 + 2t_2$$

It's better to put **shorter** jobs first.

# Algorithm

Guess: sort all jobs according to  $t_i/w_i$ :

$$\frac{t_{i_1}}{w_{i_1}} \leq \frac{t_{i_2}}{w_{i_2}} \leq \dots \leq \frac{t_{i_n}}{w_{i_n}}$$

We prove that  $i_1, \dots, i_n$  is an optimal solution.


Consider two solutions:

$$j_1, \dots, j_a, j_{a+1}, \dots, j_n \text{ and } j_1, \dots, j_{a+1}, j_a, \dots, j_n$$

that differ only by the ordering of two jobs,  $j_a$  and  $j_{a+1}$ .

# Algorithm

Denote  $x = j_a$  and  $y = j_{a+1}$



$c(b) = c'(b)$  for all  $b \notin \{x, y\}$

$c(x) = T + t_x$

$c(y) = T + t_x + t_y$

$c'(x) = T + t_x + t_y$

$c'(y) = T + t_y$


$c(x) - c'(x) = -t_y$

$c(y) - c'(y) = t_x$

$$\text{cost} - \text{cost}' = \sum_b w_b c(b) - \sum_b w_b c'(b) =$$

$$\sum_b w_b (c(b) - c'(b)) = -w_x t_y + w_y t_x$$

# Algorithm



$$\text{cost} - \text{cost}' = -w_x t_y + w_y t_x \leq 0 \text{ if and only if } \frac{t_x}{w_x} \leq \frac{t_y}{w_y} .$$

The first solution is better if and only  $\frac{t_x}{w_x} < \frac{t_y}{w_y} .$



# Algorithm

Consider an optimal solution  $i_1^*, \dots, i_n^*$ .

Is it possible that  $\frac{t_{i_a^*}}{w_{i_a^*}} > \frac{t_{i_{a+1}^*}}{w_{i_{a+1}^*}}$  for some  $a$ ?

No! Then  $i_1^*, \dots, i_n^*$  would not be an optimal solution!

# Algorithm

Consider an optimal solution  $i_1^*, \dots, i_n^*$ .

We conclude that

$$\frac{t_{i_1^*}}{w_{i_1^*}} \leq \frac{t_{i_2^*}}{w_{i_2^*}} \leq \dots \leq \frac{t_{i_n^*}}{w_{i_n^*}}$$

as required.

Are we done?

Running time?

# Summary

- Discussed greedy algorithms
- Designed algorithms for Interval/Job Scheduling and Minimum Weighted Completion Time