

Topic 2: DIMENSIONALITY REDUCTION

STAT 37710/CAAM 37710/CMSC 35400 Machine Learning
Risi Kondor, The University of Chicago

Dimensionality reduction

In ML data points are often represented as high dimensional real valued vectors




$$\mathbf{x} = (x_1, x_1, x_3, \dots, x_d)^\top \in \mathbb{R}^d.$$

The individual dimensions are called **features (attributes)**.

Example: Pixels of an image, a music file, etc.



But is the problem intrinsically high dimensional??? Often we can convert high dimensional problems to lower dimensional ones without losing too much information.

Dimensionality reduction



- Real world data often lie on or near lower dimensional structures (manifolds). (Really?)
 - Variables (features) may be correlated or dependent.
 - Physical systems have a small number of degrees of freedom (e.g., pose and lighting in Vision).
- **IDEA:** find the manifold and restrict learning algorithm to it.

Differentiable manifolds



In mathematics, a d dimensional **manifold** is a topological space such that each point has a neighborhood that is homeomorphic to \mathbb{R}^d . A differentiable manifold has additional structure, and a Riemannian manifold has a metric too \rightarrow geodesics.

Dimensionality reduction

Advantages:

- **Visualization:** humans can only imagine things in 2D or 3D.
- **Computational efficiency:** learning algorithms work faster in low dimensions.
- **Better performance:** the projection might eliminate noise.
- **Interpretability:** the vectors spanning the subspace might have interesting interpretations.

Dimensionality reduction

Dimensionality reduction is a typical **unsupervised learning** task. Two types:

- Linear:
 - Principal Component Analysis (PCA)
- Nonlinear (“manifold learning”):
 - Multidimensional scaling
 - Locally linear embedding
 - Isomap
 - Laplacian Eigenmaps
 - Stochastic neighbor embedding
 - etc.

Fact 1

If a matrix $A \in \mathbb{R}^{d \times d}$ is symmetric, then its (normalized) eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_d$ form an orthonormal basis for \mathbb{R}^d .

Note: If the eigenvalues are not distinct, then the eigenvectors are not unique. However, there is always some choice of eigenvectors which forms an orthonormal basis.

Fact 2 (Rayleigh quotient)

Let $\mathbf{v}_1, \dots, \mathbf{v}_d$ be the normalized eigenvectors of a symmetric matrix $A \in \mathbb{R}^{d \times d}$ and let $\lambda_1 < \lambda_2 < \dots < \lambda_d$ be the corresponding eigenvalues. Then


$$\operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d \setminus \{0\}} \frac{\mathbf{w}^\top A \mathbf{w}}{\|\mathbf{w}\|^2} = \mathbf{v}_1.$$

Similarly,


$$\operatorname{argmax}_{\mathbf{w} \in \mathbb{R}^d \setminus \{0\}} \frac{\mathbf{w}^\top A \mathbf{w}}{\|\mathbf{w}\|^2} = \mathbf{v}_d.$$

Principal Component Analysis

The principal directions in data



Finding the principal subspace



How can we find the most relevant subspace for the data? By finding a basis for it. The individual basis vectors are called the **principal components**.

The first principal component

Given a data set $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ of n vectors in \mathbb{R}^d , what is the direction that is most informative for this data?

1. First center the data: $\mathbf{x}_i \leftarrow \mathbf{x}_i - \boldsymbol{\mu}$ where $\boldsymbol{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$.
2. Find the unit vector \mathbf{p}_1 that is the solution to

$$\mathbf{p}_1 = \arg \max_{\|\mathbf{v}\|=1} \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i \cdot \mathbf{v})^2. \quad (1)$$

This vector is called the first **principal component** of the data.

The first principal component

Theorem. The first principal component, \mathbf{p}_1 , is the eigenvector \mathbf{v}_d of the **sample covariance matrix**

$$\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top$$

with largest eigenvalue.

Proof.

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i \cdot \mathbf{v})^2 &= \frac{1}{n} \sum_{i=1}^n (\mathbf{v}^\top \mathbf{x}_i)(\mathbf{x}_i^\top \mathbf{v}) = \frac{1}{n} \sum_{i=1}^n \mathbf{v}^\top (\mathbf{x}_i \mathbf{x}_i^\top) \mathbf{v} = \\ &= \frac{1}{n} \mathbf{v}^\top \left(\sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top \right) \mathbf{v} = \mathbf{v}^\top \hat{\Sigma} \mathbf{v}. \end{aligned}$$

Since $\|\mathbf{v}\| = 1$, (1) is equivalent to the Rayleigh quotient optimization problem

$$\mathbf{p}_1 = \arg \max_{\mathbf{v} \in \mathbb{R}^d \setminus \{0\}} \frac{\mathbf{v}^\top \hat{\Sigma} \mathbf{v}}{\|\mathbf{v}\|},$$

so \mathbf{p}_1 is indeed the eigenvector \mathbf{v}_d of A with largest eigenvalue.

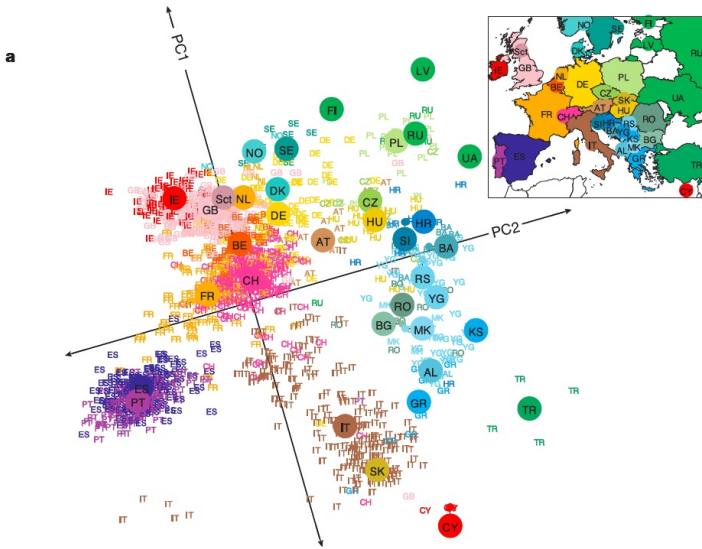
Further principal components

Recall that $\hat{\Sigma}$ can be written as

$$\hat{\Sigma} = \sum_{i=1}^d \lambda_i \mathbf{v}_i \mathbf{v}_i^{\top}.$$

After we've found the first principal component $\mathbf{p}_1 = \mathbf{v}_d$, project the data to $\text{span}\{\mathbf{v}_1, \dots, \mathbf{v}_{d-1}\}$. This just removes $\lambda_d \mathbf{v}_d \mathbf{v}_d^{\top}$ from the sum. So the second principal component is $\mathbf{p}_2 = \mathbf{v}_{d-1}$, and so on.

DNA data



[Matthew Stephens, John Novembre]

Eigenfaces




[Christopher de Cora]

Reconstruction from eigenfaces




[Christopher de Cora]

Example: digits



These are the EVectors for the four largest EValues.

- Often the eigenvalues drop off rapidly (e.g., exponentially)
- Sometimes there is a sharp drop somewhere, called the **spectral gap** → natural place to put cut-off



[Source: Peter Orbanz]

Summary of PCA

Advantages:



- Finds best projection
- Rotationally invariant

Disadvantages:


- Full PCA is expensive to compute
- Components not sparse
- Sensitive to outliers
- Linear

NONLINEAR DIMENSIONALITY REDUCTION

- If the data lies close to a linear subspace of \mathbb{R}^d , PCA can find it.
- But what if the data lies on a nonlinear **manifold**? Data which at first looks very high dimensional often really has low dimensional structure.



General principle



$$d_X(x_i, x_j) \approx |z_i - z_j|$$

Find a map $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^p$ that maps the manifold to a lower dimensional Euclidean space in a way that preserves local distances as much as possible (some methods can only map individual data points not the whole of \mathbb{R}^d).

Question: Can this always be done? Depends on the topology.

Methods

- Multidimensional Scaling
- Isomap
- Locally Linear Embedding
- Laplacian Eigenmaps
- SNE, etc..

Multidimensional scaling (MDS)

Classical MDS

- **Input:** n data points $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$.
- **Output:** n corresponding lower dimensional points $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^p$ (with $p \ll d$) that minimize the so-called *strain*

$$\mathcal{E}_{\text{CMDS}} = \|D - D^*\|_{\text{Frob}}^2 = \sum_{i,j} (D_{i,j} - D_{i,j}^*)^2,$$


where $D_{i,j} = \|\mathbf{x}_i - \mathbf{x}_j\|^2$ and $D_{i,j}^* = \|\mathbf{y}_i - \mathbf{y}_j\|^2$.

The Gram matrix

The **Gram matrix** of $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ is the $n \times n$ positive semidefinite matrix

$$G_{i,j} = \mathbf{x}_i \cdot \mathbf{x}_j.$$

(Again, we assume that the data has been centered, i.e., $\sum_i \mathbf{x}_i = 0$.)



Jørgen Pedersen Gram
1850–1916

Exercise: Prove that if $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$, then $\text{rank}(G) \leq d$.

Classical MDS

Proposition 1. The CMDS problem can equivalently be written as minimizing

$$\mathcal{E} = \| G - G^* \|_{\text{Frob}}^2,$$

where G is the centered Gram matrix of $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ and G^* is the Gram matrix of $\{\mathbf{y}_1, \dots, \mathbf{y}_n\}$.

Approach:

1. Compute the centered Gram matrix G .
2. Solve $G^* = \operatorname{argmin}_{\tilde{G} \succeq 0, \operatorname{rank}(\tilde{G}) \leq p} \| \tilde{G} - G \|_{\text{Frob}}^2$.
3. Find $\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n \in \mathbb{R}^p$ with Gram matrix G^* .

Classical MDS

Proposition 2. Let $G = Q\Lambda Q^\top$ be the eigendecomposition of the Gram matrix with $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$ and $\lambda_1 \geq \dots \geq \lambda_d$. Then

$$\operatorname{argmin}_{\tilde{G} \succeq 0, \operatorname{rank}(\tilde{G}) \leq p} \|\tilde{G} - G\|_{\text{Frob}}^2 = Q\Lambda^*Q^\top,$$

where $\Lambda^* = \text{diag}(\lambda_1, \dots, \lambda_p, 0, 0, \dots)$.

Exercise: Prove this proposition.

Gram \rightarrow Data

Proposition 3. Let $G \in \mathbb{R}^{n \times n}$ be a p.s.d. matrix of rank d with eigen-decomposition

$$G = Q\Lambda Q^\top.$$

Let $\mathbf{x}_i = [Q\Lambda^{1/2}]_{i,*}^\top$. Then the Gram matrix of $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ is G .

Notation:

- $M_{i,*}$ denotes the i 'th row of M .
- Given $D = \text{diag}(d_1, \dots, d_m)$, $D^p := \text{diag}(d_1^p, \dots, d_m^p)$.

Exercise: Prove this proposition.


Summary of Classical MDS

1. Compute the centered Gram matrix G (see homework for how).
2. Compute the eigendecomposition $Q\Lambda Q^\top$ of G .
3. Assuming $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$ and $\lambda_1 \geq \dots \geq \lambda_d$, set $\Lambda^* = \text{diag}(\lambda_1, \dots, \lambda_p, 0, 0, \dots)$ and $G^* = Q\Lambda^*Q^\top$.
4. Let $\mathbf{y}_i = [Q\Lambda^{1/2}]_{i,*}^\top$.

Isomap

Tenenbaum, de Silva & Langford, 2000

Isomap



1. Convert data into a graph (e.g., a symmetrized k -nn graph).
2. Compute all pairs shortest path distances.
3. Use MDS to compute $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^p$ that tries to preserve these distances.

Underlying assumptions:

1. Data lies on a manifold.
2. Goedsic distance on manifold is approximated by distance in the graph.
3. The optimal embedding preserves these distances as much as possible.

Shortest path distances

Let \mathcal{G} be a weighted graph with vertex set $\{1, 2, \dots, n\}$, and a distance $(\delta_{i,j})_{i,j=1}^n$ on each edge. If i and j are not neighbors, then set $\delta_{i,j} = \infty$. If $i = j$, then set $\delta_{i,j} = 0$.

The shortest path distance in \mathcal{G} from i to j is

$$d(i, j) = \min_{(v_1, v_2, \dots, v_\ell) \in \mathcal{P}(i, j)} \sum_{k=1}^{\ell-1} \delta_{v_k, v_{k+1}},$$

where \mathcal{P} is the set of paths that start at i and end at j (i.e., $v_1 = i$ and $v_\ell = j$).

Shortest path distances

Proposition. The matrix D of all pairwise distances ($D_{i,j} = d(i,j)$) can be computed in $O(n^3)$ time.

Proposition. Let $D^{(k)}$ be the matrix of shortest path distances along the restricted set of paths where each intermediate vertex comes from $\{1, 2, \dots, k\}$. Then $D^{(k)}$ can be computed from $D^{(k-1)}$ in $O(n^2)$ time.


Floyd–Warshall algorithm

```
INPUT: matrix  $A$  with  $A_{i,j} = \delta_{i,j}$  as on previous slide;  
for  $k = 1$  to  $n$  {  
  for  $i = 1$  to  $n$  {  
    for  $j = 1$  to  $n$  {  
      if  $(A_{i,j} > A_{i,k} + A_{k,j})$  then  $A_{i,j} \leftarrow A_{i,k} + A_{k,j}$  ;  
    }  
  }  
}
```

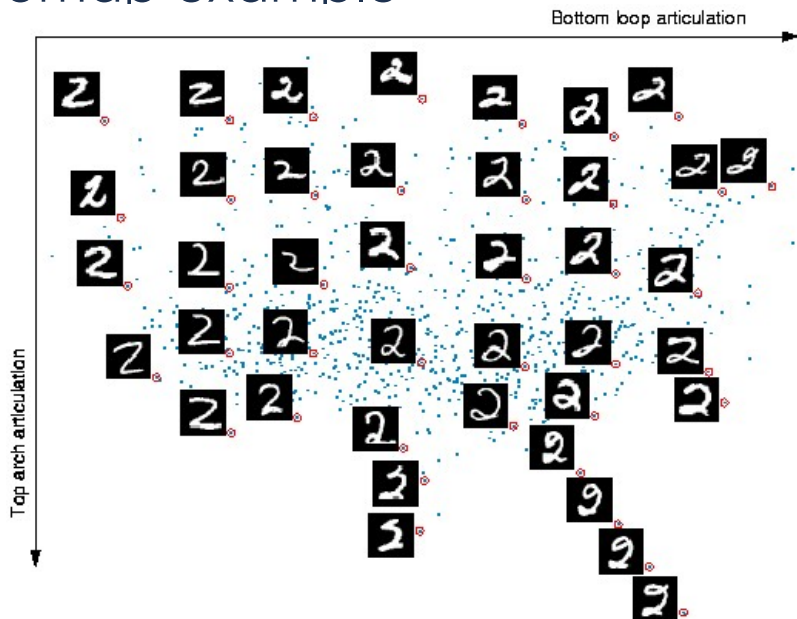
OUTPUT: matrix A , in which $A_{i,j}$ is shortest path distance from vertex i to j

Overall complexity: $O(n^3)$.


Isomap example



Isomap example



Isomap example



Properties of Isomap

- One of the first algorithms that can deal with manifolds.
- The topology must still be that of (a patch of) \mathbb{R}^p .
- Relatively efficient computation $O(n^3)$.
- Fragile: a single mistake in k -nn graph can mess up embedding.
- Not obvious how to set k .

Locally Linear Embedding (LLE)

Roweis & Saul, 2000

LLE

Again trying to find an embedding $\mathbb{R}^D \rightarrow \mathbb{R}^d$, mapping $\mathbf{x}_i \mapsto \mathbf{y}_i$. Again start with a k -nn graph based on distances in \mathbb{R}^D .

IDEA: Each point should be approximately reconstructable as a linear combination of its neighbors (locally linear property of manifolds):

$$\mathbf{x}_i \approx \sum_{j \in \text{knn}(i)} w_{i,j} \mathbf{x}_j,$$

where $(w_{i,j})_{i,j}$ is a matrix of weights. Also have constraints $\sum_j w_{i,j} = 1$.

Now find an embedding that preserves these weights, i.e., n vectors $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^p$, such that

$$\mathbf{y}_i \approx \sum_j w_{i,j} \mathbf{y}_j$$

for the same matrix of weights.

Phase 1: find the weights

Do this separately for each i . Formulate it as minimizing

$$\Phi = \left\| \mathbf{x}_i - \sum_{j \in \text{knn}(i)} w_{i,j} \mathbf{x}_j \right\|^2 \quad \text{s.t.} \quad \sum_j w_{i,j} = 1.$$

Solution. Thanks to the constraint,

$$\Phi = \left\| \sum_{j \in \text{knn}(i)} w_{i,j} (\mathbf{x}_i - \mathbf{x}_j) \right\|^2 = \mathbf{w}^\top K^{(i)} \mathbf{w},$$

where $K^{(i)}$ is the local Gram matrix, $K_{j,j'}^{(i)} = (\mathbf{x}_i - \mathbf{x}_j)^\top (\mathbf{x}_i - \mathbf{x}_{j'})$, and $\mathbf{w} = (w_j)_{j \in \text{knn}(i)}$.

Phase 1: find the weights

The local optimization problem is

$$\underset{\mathbf{w}}{\text{minimize}} \quad \mathbf{w}^\top K^{(i)} \mathbf{w} \quad \text{s.t.} \quad \mathbf{w}^\top \mathbf{1} = 1.$$

Introduce the Lagrangian:

$$\mathcal{L}(\lambda) = \mathbf{w}^\top K^{(i)} \mathbf{w} - \lambda(\mathbf{w}^\top \mathbf{1} - 1)$$

and solve

$$\frac{\partial}{\partial w_j} \mathcal{L}(\mathbf{w}) = [2K^{(i)} \mathbf{w} - \lambda \mathbf{1}]_j = 0 \quad j \in \text{knn}(i)$$

$$\mathbf{w} = \lambda (K^{(i)})^{-1} \mathbf{1} \quad \text{enforcing constraints:} \quad \mathbf{w} = \frac{(K^{(i)})^{-1} \mathbf{1}}{\| (K^{(i)})^{-1} \mathbf{1} \|_1}.$$


Phase 2: find the \mathbf{y}_i 's

Now minimize (w.r.t. $\mathbf{y}_1, \dots, \mathbf{y}_n$)

$$\Psi = \sum_i \left\| \mathbf{y}_i - \sum_j w_{i,j} \mathbf{y}_j \right\|^2 \quad s.t. \quad \sum_i \mathbf{y}_i = 0 \quad \frac{1}{n} \sum_i \mathbf{y}_i \mathbf{y}_i^\top = I.$$

Solution.


$$\Psi = \sum_{i,j} \mathbf{y}_i^\top M \mathbf{y}_j \dots$$



Laplacian Eigenmaps

Belkin and Niyogi, 2002

Spectral Graph Theory



Spectral graph theory is about relating functions on graphs (i.e., $f: V \rightarrow \mathbb{R}$ where V is the vertex set of the graph) to the structure of the graph.

Unweighted graphs

Let \mathcal{G} be an unweighted, undirected graph with vertex set $V = \{1, 2, \dots, n\}$ and edge set $E \subseteq V \times V$.

- The **adjacency matrix** of \mathcal{G} is the matrix $A \in \{0, 1\}^{n \times n}$ with

$$A = \begin{cases} 1 & \text{if } i \sim j \\ 0 & \text{otherwise,} \end{cases}$$

where $i \sim j$ means that vertices i and j are adjacent.

- The **degree matrix** of \mathcal{G} is $D = \text{diag}(d(1), d(2), \dots, d(n))$, where $d(i)$ is the degree (number of neighbors) of vertex i .
- The **Laplacian matrix** of \mathcal{G} is

$$L = D - A.$$

Laplacian as a quadratic form

The Laplacian can be written as

$$L = \sum_{i \sim j} E_{i,j} \quad \text{where} \quad [E_{i,j}]_{p,q} = \begin{cases} 1 & \text{if } p = q = i \text{ or } p = q = j \\ -1 & \text{if } (p,q) = (i,j) \text{ or } (p,q) = (j,i) \\ 0 & \text{otherwise.} \end{cases}$$

Therefore we have the fundamental identity that for any $f \in \mathbb{R}^n$,

$$f^\top L f = \sum_{i \sim j} (f(i) - f(j))^2.$$

Equivalently (and confusingly),

$$f^\top L f = \frac{1}{2} \sum_{(i,j) \in E} (f(i) - f(j))^2.$$

Exercise: Prove that L is a psd matrix.

Weighted graphs

Let \mathcal{G} be a weighted, undirected graph with edge weights $(w_{i,j})_{i,j}$. Note that $w_{i,j}=w_{j,i}$, and if $i \not\sim j$, then $w_{i,j}=0$.

- The **adjacency matrix** of \mathcal{G} is the matrix $A \in (\mathbb{R}^+)^{n \times n}$ with

$$A = \begin{cases} w_{i,j} & \text{if } i \neq j \\ 0 & \text{if } i = j. \end{cases}$$

- The **degree matrix** of \mathcal{G} is $D = \text{diag}(d(1), d(2), \dots, d(n))$, where

$$d(i) = \sum_{j \neq i} w_{i,j}.$$

- The **Laplacian matrix** of \mathcal{G} is


$$L = D - A.$$

The normalized Laplacian

When the degree distribution is uneven, it is often much better to work with the **normalized Laplacian**

$$\tilde{L} = D^{-1/2} L D^{-1/2} = I - D^{-1/2} A D^{-1/2}.$$

Example: cycle graph



$$f_k(v_i) = \sin(2\pi ki/n)$$


$$k = 1, 2, \dots, \lfloor n/2 \rfloor$$

$$g_k(v_i) = \cos(2\pi ki/n)$$

$$k = 0, 1, 2, \dots, \lfloor (n-1)/2 \rfloor,$$

Example: path graph

Eigenvectors of path graph



Connectivity

Theorem

The multiplicity of 0 in the spectrum of L (i.e., the number of zero eigenvalues) is the number of connected components of \mathcal{G} .

Fiedler vector

Let $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$ be the eigenvalues of L , and v_1, v_2, \dots, v_n be the corresponding normalized eigenvectors.

- By the above, $\lambda_1 = 0$ and $v_1 = \frac{1}{\sqrt{n}} \mathbf{1}$ for any graph.
- The second eigenvector v_2 , is called the **Fiedler vector**, and is particularly informative about how to cluster the graph.

Cheeger's inequality

Let $S \subset V$, $\bar{S} = V \setminus S$, and $E(S, \bar{S}) = \sum_{i \in S} \sum_{j \in \bar{S}} w_{i,j}$. Further for any $W \subseteq V$, let $d(W) = \sum_{i \in W} d(i)$.

- The **conductance** of S is defined as


$$\phi(S) = d(V) \frac{E(S, \bar{S})}{d(S) d(\bar{S})}.$$

- The conductance of the whole graph is $\phi_G = \min_{S \subset V} \phi(S)$.
- Cheeger's inequality** states that



$$\frac{\phi_G^2}{2d_{\max}} \leq \lambda_2 \leq \phi_G,$$

where d_{\max} is the maximum degree of any vertex in G .

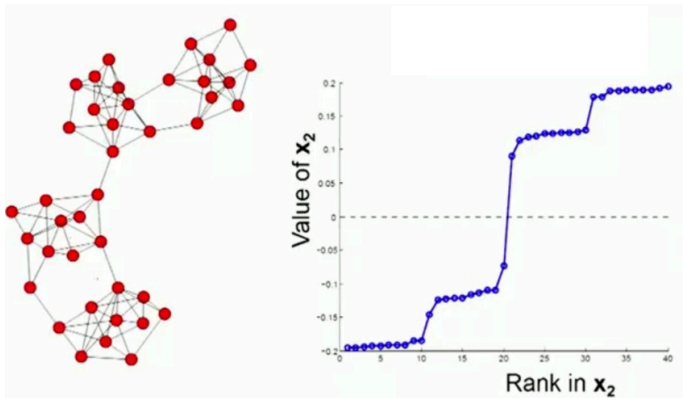
Example



Example



Example



The first few eigenvectors can be used for clustering → spectral graph partitioning

The Laplace–Beltrami operator

The graph Laplacian is the discrete analog of the Laplace–Beltrami operator.

- The Laplacian operator on \mathbb{R}^d is

$$\Delta = \nabla^2 = \frac{\partial^2}{\partial x_1^2} + \frac{\partial^2}{\partial x_2^2} + \dots + \frac{\partial^2}{\partial x_d^2}.$$

- More generally, the Laplace–Beltrami operator on a d dimensional Riemannian manifold with metric tensor g is

$$\Delta = \frac{1}{\sqrt{\det g}} \sum_{i,j=1}^d \partial_i \sqrt{\det g} g^{i,j} \partial_j.$$

The graph Laplacian can be regarded as a discretization of these operators.

Discretization of Laplacian

- In \mathbb{R} , the (finite difference) discretization of $\nabla = \frac{\partial}{\partial x}$ is derived from

$$(\nabla f)(x) = \left(\frac{\partial}{\partial x} f \right)(x) = \frac{f(x+h/2) - f(x-h/2)}{h}.$$

- The discretization of Δ is derived from

$$\begin{aligned} (\Delta f)(x) &= (\nabla(\nabla f))(x) = \frac{(\nabla f)(x+h/2) - (\nabla f)(x-h/2)}{h} = \\ &= \frac{f(x-h) - 2f(x) + f(x+h)}{h^2}. \end{aligned}$$

If we regard f as a vector, $\mathbf{f} = (\dots, f(x-h), f(x), f(x+h), \dots)^\top$, then the latter is just $-L\mathbf{f}/h^2$, where L is the Laplacian of the line graph. Similarly for grids on \mathbb{R}^d . $\langle f, \Delta f \rangle$ is a natural measure of roughness of f \rightarrow sheds new light on L as a quadratic form.

The heat equation

The flow of heat in a homogenous medium is governed by the equation

$$\frac{\partial}{\partial t} f(\mathbf{x}, t) = \kappa \Delta f(\mathbf{x}, t).$$

Δ is a negative definite self-adjoint operator. Solution to this is

$$f(\mathbf{x}, t) = e^{\kappa t \Delta} f(\mathbf{x}, 0) \quad \text{where} \quad e^T := I + T + \frac{1}{2}T^2 + \frac{1}{6}T^3 + \dots$$


In particular, if our domain \mathcal{M} is compact, then the eigenfunctions of Δ , i.e., $\Delta g_i = \lambda_i g_i$ form a basis for \mathcal{M} and

$$f(\mathbf{x}, 0) = \sum_i \alpha_i g_i \qquad f(\mathbf{x}, t) = \sum_i e^{\lambda_i \kappa t} \alpha_i g_i.$$

The long time behavior of the system is determined by the low $|\lambda_i|$ modes!!!

Laplacian Eigenmaps

[Belkin&Niyogi]



- Turn dimensionality reduction into a graph problem by forming knn-mesh, possibly weighted by $w_{i,j} = \exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / (2\sigma^2))$
- Embed according to first p non-zero e-value e-vectors:

$$\phi: V \rightarrow \mathbb{R}^p \quad i \mapsto \begin{pmatrix} v_1(i) \\ \vdots \\ v_{p+1}(i) \end{pmatrix}$$

- Intuition: these are the smoothest functions on the graph, and they give global coordinates

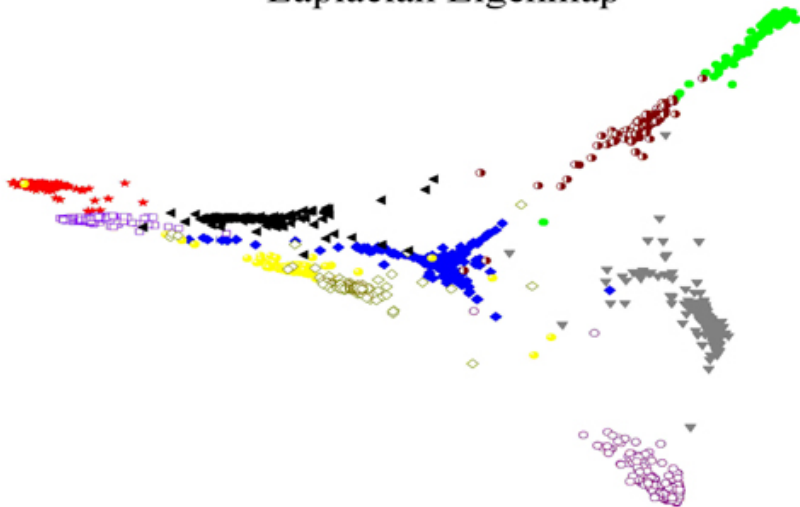
Laplacian Eigenmaps: detail

Formulate the problem as minimizing the strain


$$\mathcal{E} = \sum_{i,j} w_{i,j} \|\mathbf{y}_i - \mathbf{y}_j\|^2 = 2\text{tr}(Y^\top LY).$$

Adding the additional constraint $Y^\top DY = I$, after some algebra, this leads to the generalized eigenvalue problem $LY = DY$.


Laplacian Eigenmap




Three different metrics



Euclidean




Geodesic




Diffusion

Laplacian eigenmaps corresponds to PCA w.r.t. the diffusion metric on the manifold, because the diffusion (heat) kernel is exactly $e^{-\beta L}$ [K and Lafferty, 2001].


$(n,m)=(1,1), \lambda=0.351$




$(n,m)=(1,2), \lambda=0.648$




$(n,m)=(1,3), \lambda=0.955$




$(n,m)=(2,1), \lambda=0.444$




$(n,m)=(2,2), \lambda=0.702$




$(n,m)=(2,3), \lambda=0.993$




$(n,m)=(3,1), \lambda=0.566$




$(n,m)=(3,2), \lambda=0.785$



$(n,m)=(3,3), \lambda=1.05$




Original data, N=2000




MDS, corr=0.7938




Isomap, corr=0.9999




LLE, k= 20, corr=0.5286




HessianLLE, k= 20, corr=0.9003




LTSA, k= 20, corr=0.9003



KernelPCA, poly, corr=0.4236



DiffusionMaps, corr=0.7022



AutoEncoderRBM, corr=0.5645

